

Question 2 (Henry)

2026-10-09

Do students who have taken both 131 and 134 report greater interest in machine-learning-related fields than students who have taken 131 only?

Our second question surrounds the association between taking certain PSTAT courses and having specific areas of interest. Specifically, we want to see if taking PSTAT 131 and/or PSTAT 134 has any correlation in machine learning domain interests. We chose the 131/134 courses because its course description fits the topic of machine learning and its techniques pretty well. It is a course that is well known in the PSTAT department, and with the ongoing advancements in artificial intelligence and machine learning, we only feel that the latter is worth exploring through our survey data.

```
problem2 <- read_csv("~/Desktop/PSTAT197A/module0-class-survey-team9/data/processed/q2/problem2.csv")
```

```
## Rows: 49 Columns: 15
## -- Column specification -----
## Delimiter: ","
## chr (11): prog.prof, math.prof, stat.prof, upper_div_courses, courses, domai...
## dbl (4): response_id, prog.comf, math.comf, stat.comf
##
## i Use 'spec()' to retrieve the full column specification for this data.
## i Specify the column types or set 'show_col_types = FALSE' to quiet this message.
```

```
dim(problem2)
```

```
## [1] 49 15
```

```
head(problem2)
```

```
## # A tibble: 6 x 15
##   response_id prog.prof prog.comf math.prof math.comf stat.prof stat.comf
##   <dbl> <chr>         <dbl> <chr>         <dbl> <chr>         <dbl>
## 1      1001 int             4 int             4 adv             5
## 2      1002 int             4 int             3 int             5
## 3      1003 beg             3 int             3 adv             4
## 4      1004 adv             4 adv             5 adv             5
## 5      1005 adv             4 int             4 adv             5
## 6      1006 int             4 int             2 adv             4
## # i 8 more variables: upper_div_courses <chr>, courses <chr>,
## #   domain_specialization <chr>, research <chr>, type <chr>, language <chr>,
## #   domains <chr>, areas <chr>
```

We have already joined the two datasets (**background-clean.csv** and **interest-clean.csv**) by `response_id`. We can see that there are 49 responders who consented to sharing both responses. While there are 15 columns, we will mainly be using `courses` and `areas`.

Data Validation

We are assuming that students who have taken PSTAT 134 have also taken PSTAT 131 since it is a pre-requisite to that course. We checked this assumption by ensuring that every student who listed PSTAT 134 also included PSTAT 131.

```
course_check <- problem2 %>%
  mutate(
    took_131 = str_detect(courses, "\\bPSTAT 131\\b"),
    took_134 = str_detect(courses, "\\bPSTAT 134\\b")
  )

# Show everyone who took PSTAT 134 but not PSTAT 131
course_check %>%
  filter(took_134 & !took_131) %>%
  select(response_id, courses)
```

```
## # A tibble: 0 x 2
## # i 2 variables: response_id <dbl>, courses <chr>
```

Since there are no students who took PSTAT 134, but not PSTAT 131, our data is valid.

Now let's see the potential options for area interests:

```
possible_values <- unique(unlist(strsplit(as.character(problem2$areas), split = ",")))
possible_values
```

```
## [1] "Spatial statistics or time series analysis"
## [2] " Data acquisition and engineering"
## [3] " Data visualization and interactive dashboards"
## [4] " Predictive modeling"
## [5] " generally"
## [6] " Statistical models and inference"
## [7] "Analysis or classification of images"
## [8] " Spatial statistics or time series analysis"
## [9] " Natural language processing and analysis of text"
## [10] " Model deployment and software or web integrations"
## [11] " Databases"
## [12] " Algorithms"
## [13] " Deep learning and neural networks"
## [14] "Deep learning and neural networks"
## [15] "Data acquisition and engineering"
## [16] "Model deployment and software or web integrations"
```

Out of the eleven options (not counting spacing discrepancies), we will define 'machine learning' interests as one of the following six options:

```
ml_areas <- c(
  "Analysis or classification of images",
  "Deep learning and neural networks",
  "Data acquisition and engineering",
  "Natural language processing and analysis of text",
  "Predictive modeling, generally",
  "Statistical models and inference, generally"
)
```

Again, this part is subjective since most of the options can be interpreted to fit into the scope of machine learning. However, we only chose six out of the eleven topics to provide an adequate representation of what is closest to machine learning.

PSTAT 131 vs non-PSTAT 131

To begin, we will explore students who have taken PSTAT 131 and students who have not taken PSTAT 131.

```
problem2 %>%
  mutate(
    course_group = if_else(
      str_detect(coalesce(courses, ""), "\\bPSTAT 131\\b"),
      "Took PSTAT 131",
      "Did not take PSTAT 131"
    )
  ) %>%
  count(course_group, name = "Number of Students") %>%
  mutate(
    Proportion = `Number of Students` / sum(`Number of Students`)
  ) %>%
  knitr::kable(
    col.names = c(
      "PSTAT 131 Enrollment",
      "Number of Students",
      "Proportion"
    ),
    digits = 3,
    caption = "Students by PSTAT 131 Enrollment"
  )
```

Table 1: Students by PSTAT 131 Enrollment

PSTAT 131 Enrollment	Number of Students	Proportion
Did not take PSTAT 131	5	0.102
Took PSTAT 131	44	0.898

Here, we can see that only 5 students have not taken PSTAT 131, totaling around 10% of the respondents. On the other hand, 90% of students have taken PSTAT 131 at 44 students.

```

ml_data <- problem2 %>%
  mutate(
    took_131 = str_detect(
      coalesce(courses, ""),
      "\\bPSTAT 131\\b"
    ),
    areas = coalesce(areas, "")
  ) %>%
  separate_rows(areas, sep = ",\\s*") %>%
  mutate(selected = areas %in% ml_areas) %>%
  group_by(response_id, took_131) %>%
  summarise(
    n_ml_areas = sum(selected),
    .groups = "drop"
  ) %>%
  mutate(
    course_group = if_else(
      took_131,
      "Took PSTAT 131",
      "Did not take PSTAT 131"
    )
  )
ml_data

```

```

## # A tibble: 49 x 4
##   response_id took_131 n_ml_areas course_group
##         <dbl> <lgl>         <int> <chr>
## 1         1001 TRUE             1 Took PSTAT 131
## 2         1002 FALSE            3 Did not take PSTAT 131
## 3         1003 TRUE             4 Took PSTAT 131
## 4         1004 TRUE             2 Took PSTAT 131
## 5         1005 TRUE             4 Took PSTAT 131
## 6         1006 TRUE             2 Took PSTAT 131
## 7         1007 TRUE             4 Took PSTAT 131
## 8         1008 TRUE             3 Took PSTAT 131
## 9         1009 TRUE             4 Took PSTAT 131
## 10        1010 TRUE             3 Took PSTAT 131
## # i 39 more rows

```

```

ml_data %>%
  group_by(course_group) %>%
  summarise(
    n = n(),
    mean_ml_areas = mean(n_ml_areas),
    median_ml_areas = median(n_ml_areas),
    .groups = "drop"
  )

```

```


## # A tibble: 2 x 4
##   course_group          n mean_ml_areas median_ml_areas
##   <chr>          <int>         <dbl>         <dbl>
## 1 Did not take PSTAT 131     5           2           3

```

```
## 2 Took PSTAT 131          44          2.5          2.5
```

Here we have a contingency

```
ggplot(ml_data, aes(x = course_group, y = n_ml_areas,  
                    fill = course_group)) +  
  geom_boxplot() +  
  labs(  
    title = "ML-Related Interests by PSTAT 131 Enrollment",  
    x = "Course Group",  
    y = "Number of ML-Related Areas Selected"  
  ) +  
  theme_minimal() +  
  theme(legend.position = "none")
```



```
wilcox.test(  
  n_ml_areas ~ course_group,  
  data = ml_data,  
  exact = FALSE  
)
```

```
##  
## Wilcoxon rank sum test with continuity correction  
##  
## data:  n_ml_areas by course_group
```

```
## W = 91, p-value = 0.5252
## alternative hypothesis: true location shift is not equal to 0
```


Both PSTAT 131 and PSTAT 134 vs Only PSTAT 131

```
ml_data <- problem2 %>%
  mutate(
    took_131 = str_detect(courses, "\\bPSTAT 131\\b"),
    took_134 = str_detect(courses, "\\bPSTAT 134\\b"),
    course_group = case_when(
      took_131 & took_134 ~ "Both courses",
      took_131 & !took_134 ~ "PSTAT 131 only",
      TRUE ~ NA_character_
    ),
    areas = coalesce(areas, "")
  ) %>%
  filter(!is.na(course_group)) %>%
  separate_rows(areas, sep = ",\\s*") %>%
  mutate(selected = areas %in% ml_areas) %>%
  group_by(response_id, course_group) %>%
  summarise(
    n_ml_areas = sum(selected),
    .groups = "drop"
  )
```

```
ml_data %>%
  group_by(course_group) %>%
  summarise(
    n = n(),
    mean_domains = mean(n_ml_areas),
    median_domains = median(n_ml_areas),
    .groups = "drop"
  )
```

```
## # A tibble: 2 x 4
##   course_group      n mean_domains median_domains
##   <chr>          <int>         <dbl>         <int>
## 1 Both courses    13          2.54            3
## 2 PSTAT 131 only  31          2.48            2
```

```
ggplot(ml_data, aes(x = course_group, y = n_ml_areas,
                    fill = course_group)) +
  geom_boxplot() +
  labs(
    title = "ML-Related Interests by Course-Taking Group",
    x = "Course-Taking Group",
    y = "Number of ML-Related Areas Selected"
  ) +
  theme_minimal() +
  theme(legend.position = "none")
```



Drawbacks

Some drawbacks about this research question is that the survey allowed the respondent as many choices they wanted for area interests. In other words, someone could have checked all the boxes and make the cohorts even harder to determine. Ideally, we would have each person only selected one topic they are most interested in.